

ZURU™

GAN - Theory and Applications

Emanuele Ghelfi

Paolo Galeone

Federico Di Mattia

Michele De Simoni

April 8, 2019



PYCONX

Generative Adversarial Networks

“Adversarial Training (also called GAN for Generative Adversarial Networks) is the most interesting idea in the last 10 years of ML.”

— Yann LeCun

Generative Adversarial Networks

Two components, the **generator** and the **discriminator**:

- The **generator** G, aim is to capture the data distribution.
- The **discriminator** D, estimates the probability that a sample came from the training data rather than from G.

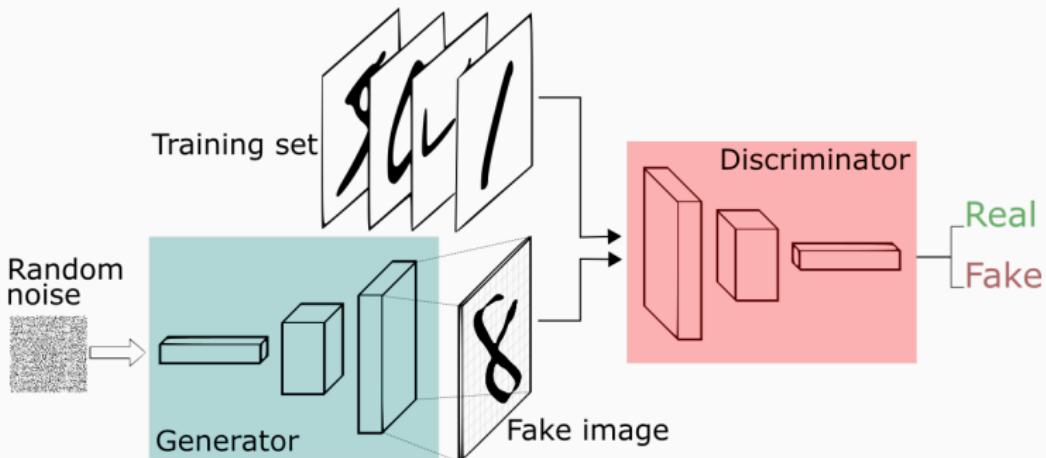


Figure 1: Credits: Silva

Generative Adversarial Networks

GANs game:

$$\min_G \max_D V_{GAN}(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (1)$$

Generative Adversarial Networks

GANs game:

$$\min_G \max_D V_{GAN}(D, G) = \underbrace{\mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]}_{\text{real samples}} \quad (1)$$

Generative Adversarial Networks

GANs game:

$$\min_G \max_D V_{GAN}(D, G) = \underbrace{\mathbb{E}_{x \sim p_{data}(x)} [\log D(x)]}_{\text{real samples}} + \underbrace{\mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]}_{\text{generated samples}} \quad (1)$$

GANs - Discriminator

- **Discriminator** needs to:

- Correctly classify **real** data:

$$\max_D \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)]. \quad (2)$$

- Correctly classify **wrong** data:

$$\max_D \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (3)$$

- The discriminator is an **adaptive loss function** that gets discarded once the generator has been trained.

GANs - Generator

- **Generator** needs to **fool** the discriminator:
 - Generate samples similar to the real one:

$$\min_G \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (4)$$

GANs - Generator

- **Generator** needs to **fool** the discriminator:

- Generate samples similar to the real one:

$$\min_G \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (4)$$

- Saturates easily Goodfellow et al. (2014).
- Change loss for generator:

$$\max_G \mathbb{E}_{z \sim p_z(z)} [\log(D(G(z)))] \quad (5)$$

GANs - Generator

- **Generator** needs to **fool** the discriminator:

- Generate samples similar to the real one:

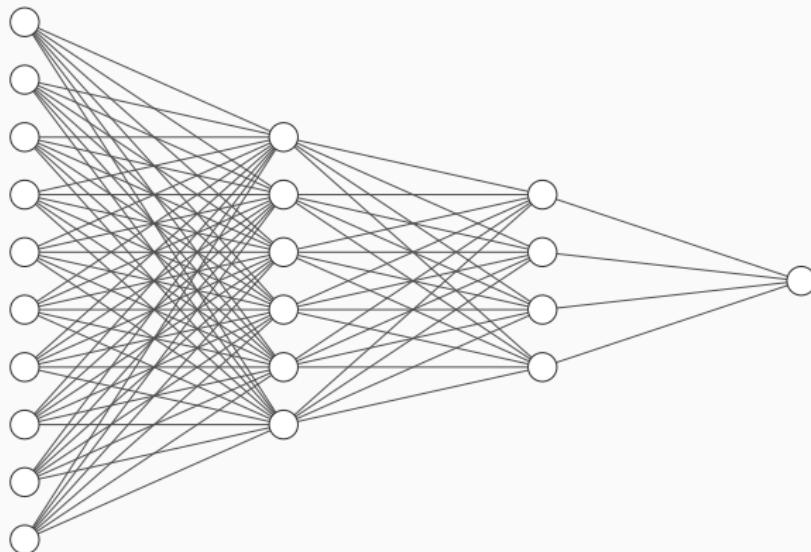
$$\min_G \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (4)$$

- Saturates easily Goodfellow et al. (2014).
- Change loss for generator:

$$\max_G \mathbb{E}_{z \sim p_z(z)} [\log(D(G(z)))] \quad (5)$$

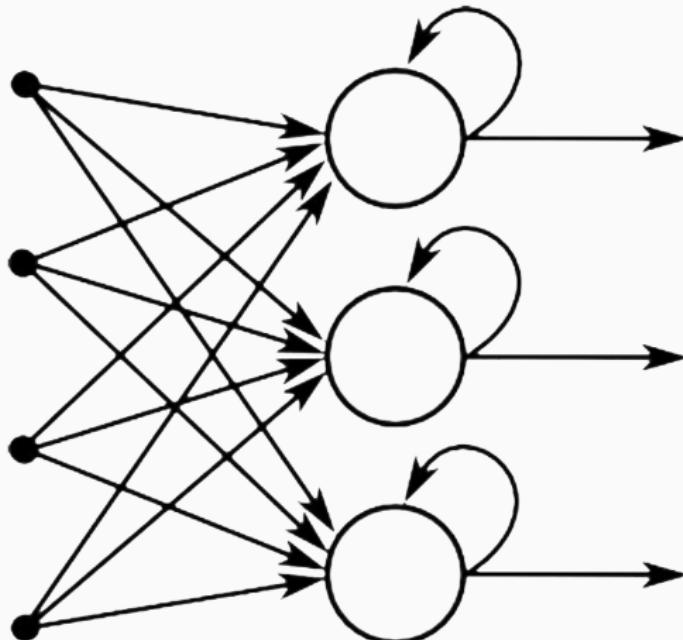
GANs - Models definition

- Different architectures for different data types.
 - Tuple of numbers? Fully Connected Neural Networks



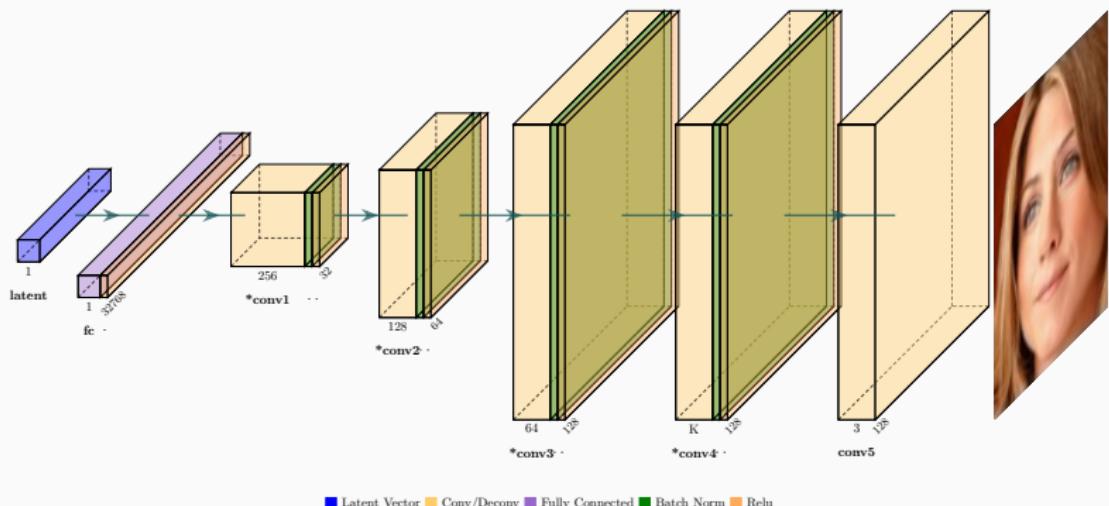
GANs - Models definition

- Different architectures for different data types.
 - Text or sequences? Recurrent Neural Networks



GANs - Models definition

- Different architectures for different data types.
 - Images? **Convolutional Neural Networks**



GANs Training

GANs - Training

- Discriminator and generator are **competing** against each other.
- **Alternating** execution of training steps.
- Use **minibatch stochastic gradient descent/ascent**.



GANs - Training - Discriminator

How to **train** the **discriminator**?

Repeat from 1 to **k**:

1. Sample minibatch of m noise samples $z^{(1)}, \dots, z^{(m)}$ from $p_g(z)$

GANs - Training - Discriminator

How to **train** the **discriminator**?

Repeat from 1 to **k**:

1. Sample minibatch of m noise samples $z^{(1)}, \dots, z^{(m)}$ from $p_g(z)$
2. Sample minibatch of m examples $x^{(1)}, \dots, x^{(m)}$ from $p_{data}(x)$

GANs - Training - Discriminator

How to **train** the **discriminator**?

Repeat from 1 to **k**:

1. Sample minibatch of m noise samples $z^{(1)}, \dots, z^{(m)}$ from $p_g(z)$
2. Sample minibatch of m examples $x^{(1)}, \dots, x^{(m)}$ from $p_{data}(x)$
3. Train the **discriminator** by stochastic gradient **ascent**:

$$\nabla_{\theta_d} \underbrace{\frac{1}{m} \sum_{i=1}^m \underbrace{\log D(x^{(i)}) + \log(1 - D(G(z^{(i)})))}_{\text{Discriminator loss}}}_{\text{Loss estimation using } m \text{ samples}}$$

GANs - Training - Generator

How to **train** the **generator**?

The update is executed **only once** and only after the turn of the discriminator is completed:

1. Sample minibatch of m noise samples $z^{(1)}, \dots, z^{(m)}$ from $p_g(z)$

GANs - Training - Generator

How to **train** the **generator**?

The update is executed **only once** and only after the turn of the discriminator is completed:

1. Sample minibatch of m noise samples $z^{(1)}, \dots, z^{(m)}$ from $p_g(z)$
2. Train the **generator** by stochastic gradient **ascent**:

$$\nabla_{\theta_g} \underbrace{\frac{1}{m} \sum_{i=1}^m \log(D(G(z^{(i))}))}_{\text{Generator loss}}$$

Loss estimation using m samples

GANs - Training - Considerations

- Optimizers: Adam, Momentum, RMSProp.
- Training phase can last for an **arbitrary number** of steps or epochs.
- Training is completed when the discriminator is **completely fooled** by the generator.
- Goal: reach a **Nash Equilibrium** where the best D can do is random guessing.

Types of GANs

Types of GANs

Two big families:

- **Unconditional** GANs (just described).
- **Conditional** GANs (Mirza and Osindero, 2014).

Conditional GANs

- Both G and D are **conditioned** on some extra information y .
- In **practice**: perform conditioning by feeding y into the discriminator and generator.

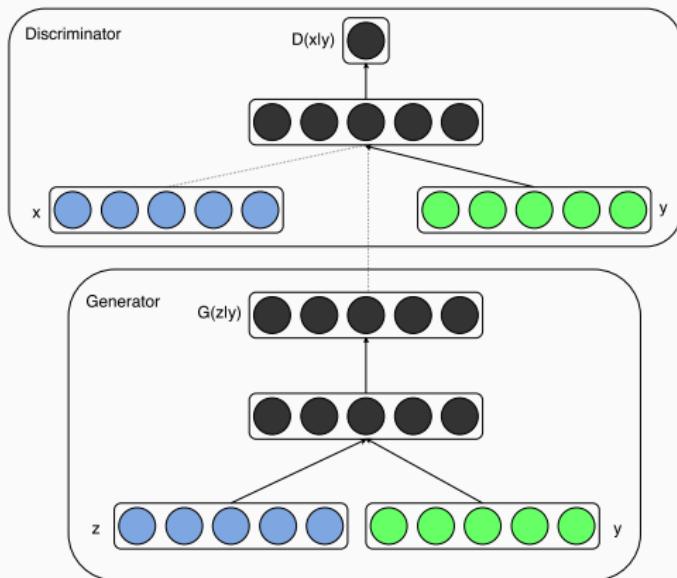


Figure 2: From Mirza and Osindero (2014)

Conditional GANs

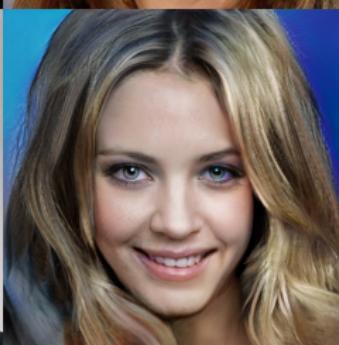
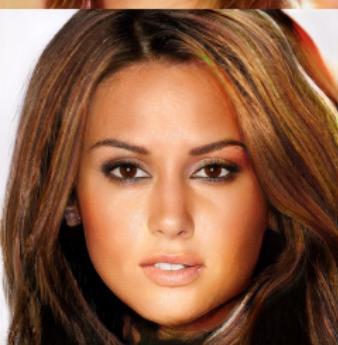
The GANs game becomes:

$$\min_G \max_D \mathbb{E}_{x \sim p_{data}(x|y)} [\log D(x, y)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z|y), y))]$$

Notice: the same representation of the condition has to be presented to both network.

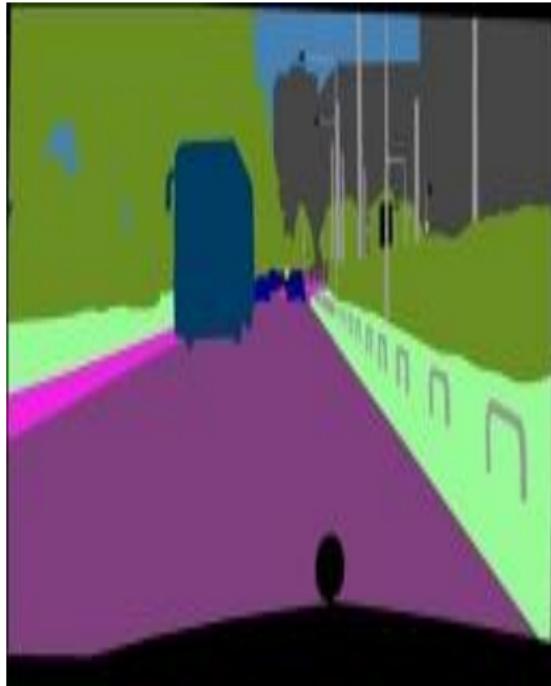
GANs Applications

Face Generation - Karras et al. (2017)



Domain Translation - Isola et al. (2016)

Input



Generated

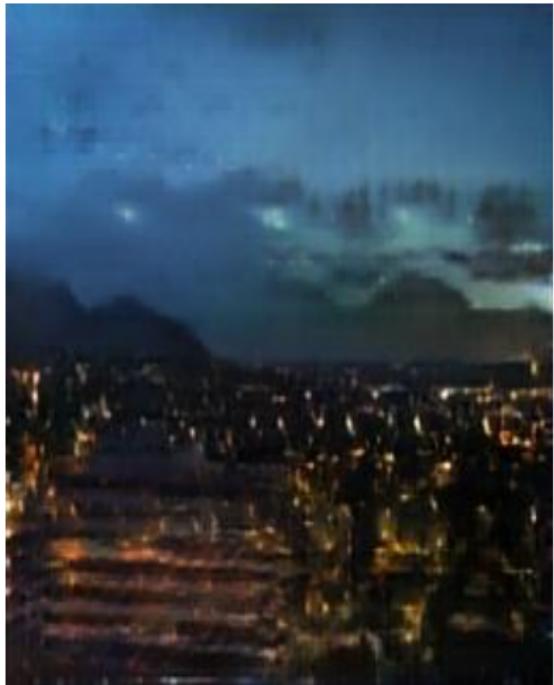


Domain Translation - Isola et al. (2016)

Input



Generated



Domain Translation - Isola et al. (2016)

Input



Generated

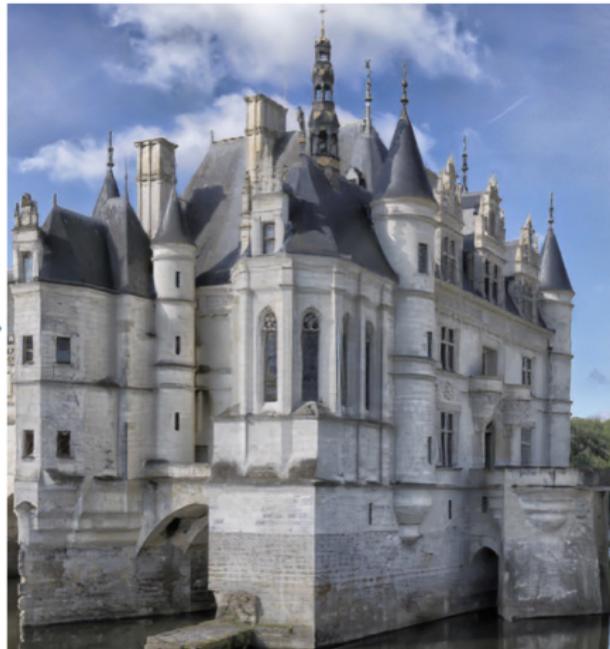


Image Super Resolution - Ledig et al. (2016)



LG Image

SRGAN



Generated Image

Thank you for your attention!
Questions?

References

- [Goodfellow et al. 2014] GOODFELLOW, Ian J. ;
POUGET-ABADIE, Jean ; MIRZA, Mehdi ; XU, Bing ;
WARDE-FARLEY, David ; OZAIR, Sherjil ; COURVILLE, Aaron ;
BENGIO, Yoshua: Generative Adversarial Networks. (2014). –
(2014)
- [Isola et al. 2016] ISOLA, Phillip ; ZHU, Jun-Yan ; ZHOU,
Tinghui ; EFROS, Alexei A.: Image-to-Image Translation with
Conditional Adversarial Networks. (2016). – (2016)
- [Karras et al. 2017] KARRAS, Tero ; AILA, Timo ; LAINE,
Samuli ; LEHTINEN, Jaakko: Progressive Growing of GANs for
Improved Quality, Stability, and Variation. (2017). – (2017)

[Ledig et al. 2016] LEDIG, Christian ; THEIS, Lucas ; HUSZAR, Ferenc ; CABALLERO, Jose ; CUNNINGHAM, Andrew ; ACOSTA, Alejandro ; AITKEN, Andrew ; TEJANI, Alykhan ; TOTZ, Johannes ; WANG, Zehan ; SHI, Wenzhe: Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network. (2016). – (2016)

[Mirza and Osindero 2014] MIRZA, Mehdi ; OSINDERO, Simon: Conditional Generative Adversarial Nets. (2014). – (2014)

[Silva] SILVA, Thalles: *An Intuitive Introduction to Generative Adversarial Networks (GANs)*